

LLM-Assisted Extraction of QSP Calibration Targets: A Validated Schema for Structured Data and Joint Bayesian Inference

Joel Eliason¹

¹Department of Biomedical Engineering, Johns Hopkins University, Baltimore, MD, USA

Abstract

Quantitative systems pharmacology (QSP) models require parameter values derived from experimental literature, yet current approaches to literature extraction face a trade-off between throughput and rigor. Parameter databases provide point estimates without uncertainty or provenance; manual expert curation yields high-quality extractions but does not scale; and large language models (LLMs) offer rapid extraction but exhibit failure modes including value hallucination and citation fabrication. We present a framework that combines LLM-assisted extraction with systematic validation to address this trade-off. The SubmodelTarget schema separates data extraction from model specification, capturing literature values with full provenance while constraining output to a machine-verifiable form. A ten-validator framework targets known LLM failure modes, including a value-in-snippet check that requires extracted numbers to appear verbatim in quoted source text. Validated targets are automatically translated to Julia/Turing.jl scripts for Bayesian inference, with automatic identification of shared parameters across targets. We demonstrate the framework on ten calibration targets for pancreatic stellate cell dynamics in pancreatic ductal adenocarcinoma (PDAC). Joint inference yields well-converged posteriors ($\hat{R} < 1.01$) for all parameters, and posterior predictive checks confirm appropriate uncertainty calibration. The framework provides a reproducible pipeline from literature to calibrated parameters while maintaining the provenance required for scientific applications.

1 Introduction

1.1 The Calibration Challenge

Quantitative systems pharmacology models translate mechanistic understanding of disease biology into predictions of drug response [Gadkar et al., 2016]. A typical QSP model comprises dozens to hundreds of ordinary differential equations representing cell populations, signaling molecules, and drug pharmacokinetics. Parameterizing these models requires values drawn from experimental literature—proliferation rates, death rates, secretion constants, binding affinities—often fifty or more parameters per model [Ribba et al., 2017]. Proper calibration with uncertainty quantification is essential not only for model credibility but also for downstream applications such as virtual population generation and virtual clinical trials, where parameter distributions directly determine predicted patient heterogeneity [Rieger et al., 2018, Craig et al., 2023].

The challenge is not finding relevant papers; it is extracting data systematically with full provenance. Each experiment reports results in different formats, with different units, time points, and uncertainty quantification. Experimental conditions vary across studies and are inconsistently documented. Converting a reported measurement into a calibration constraint requires judgment about applicability, uncertainty, and potential biases. And crucially, the reasoning behind each extraction—why this value, from this paper, interpreted this way—must be documented for reproducibility and review.

Current practice relies on manual curation into spreadsheets. This approach yields poorly documented assumptions, lost provenance, and copy-paste propagation across models. When a reviewer asks “where did this number come from?”, the answer is often difficult to reconstruct. When a new team member joins, institutional knowledge must be transferred orally. When a parameter needs updating, the original reasoning may be lost.

A further challenge arises when combining constraints from multiple experiments. Parameters often appear across different assays, and the ideal calibration would use all available data simultaneously. But without a structured representation of calibration targets, joint inference requires ad-hoc scripting for each new combination of experiments. There is no standard path from extracted data to inference code.

The rate of relevant publications outpaces what manual curation can accommodate. A scalable

approach requires automation, but automation introduces its own risks.

1.2 Isolated System Experiments

The biological literature contains a wealth of experiments that isolate specific mechanisms from the complexity of in vivo systems. In vitro proliferation assays measure cell division rates without immune infiltration or nutrient gradients. Ex vivo cytotoxicity experiments quantify killing efficiency without the confounding effects of tumor heterogeneity or drug distribution. Preclinical studies in simplified animal models can isolate specific pathway dynamics.

These isolated system experiments are valuable precisely because they remove confounding factors. When a proliferation assay reports a doubling time, that measurement directly constrains the proliferation rate parameter. The experimental design has already done the work of isolating the relevant biology. Parameters that would be unidentifiable from clinical data become identifiable from targeted in vitro measurements [Ribba et al., 2017].

The challenge is that this data is scattered across thousands of papers, reported in varying formats, with different units, time points, and uncertainty quantification. Experimental conditions (cell line, passage number, serum concentration, oxygen tension) vary across studies and are inconsistently documented. Converting a reported doubling time into a prior distribution for a rate constant requires judgment about uncertainty, applicability, and potential biases.

We use the term *calibration target* to refer to a structured representation of experimental data suitable for Bayesian inference. A calibration target captures what was measured, under what conditions, with what uncertainty, and how that measurement relates to model parameters. The goal of this work is to automate the extraction of calibration targets from literature while maintaining the provenance and rigor required for scientific applications.

1.3 Existing Approaches and Their Limitations

1.3.1 Parameter Databases

Several databases aggregate quantitative biological parameters from literature. BioNumbers provides over 13,000 entries covering cell sizes, metabolite concentrations, reaction rates, and growth constants, each linked to its source publication [Milo et al., 2010]. BRENDA contains over 5 mil-

lion data points for approximately 90,000 enzymes, including kinetic parameters manually extracted from 157,000 primary literature references [Chang et al., 2021]. Sabio-RK focuses specifically on biochemical reaction kinetics [Wittig et al., 2012]. For pharmacokinetic modeling, PK-Sim provides physiological parameters for humans and common laboratory animals as part of the Open Systems Pharmacology suite [Lippert et al., 2019].

These databases serve their intended purposes well, but they have limitations for QSP calibration. Most store point estimates without uncertainty quantification. Experimental context—the species, cell type, assay conditions, and methodological details that determine whether a value applies to a particular modeling scenario—is inconsistently documented. The databases are organized around biological entities (enzymes, metabolites, organisms) rather than the inference tasks that QSP modelers face. And critically, tracing a value back through intermediate sources to the original experimental measurement often requires substantial manual effort.

1.3.2 Biomedical Text Mining

Automated extraction of quantitative information from scientific literature has been a goal of biomedical text mining for decades. Named entity recognition can identify mentions of genes, proteins, diseases, and drugs; relation extraction can link entities into structured assertions; and numerical information extraction can pull quantities from text [Zhao et al., 2021]. Tools like PubTator and SciSpacy provide pre-trained models for biomedical NLP tasks.

The challenge for QSP parameterization is that the relevant information is often embedded in complex experimental descriptions. A reported doubling time depends on cell type, passage number, serum concentration, oxygen tension, and measurement method—details that are scattered across methods sections, figure captions, and supplementary materials. Traditional NLP pipelines struggle with this kind of distributed, context-dependent extraction. Performance on relation extraction benchmarks does not translate directly to the task of assembling coherent calibration targets from heterogeneous experimental reports [Lever et al., 2019].

1.3.3 Manual Expert Curation

Manual curation by domain experts remains the gold standard for quality. An experienced modeler can identify the relevant passages, assess methodological quality, judge applicability to the modeling

context, and document assumptions. The result is a well-reasoned parameter estimate with clear provenance.

The limitation is scale. A thorough literature review for a single QSP model can require weeks of effort. Documentation quality varies across individuals and projects. When personnel change, institutional knowledge disappears—the reasoning behind a parameter choice may be lost even if the value persists in the model. And the rate of new publications outpaces what manual curation can accommodate; PubMed adds over one million new records annually.

1.3.4 Large Language Models

Large language models offer a potential path forward. LLMs can read scientific text, follow structured templates, and extract information at scale. Recent work has explored LLM-assisted workflows for QSP model development, including semi-automated parameter extraction from literature [Saini and Farnoud, 2025]. The potential for accelerating model development is substantial.

But LLMs have well-documented failure modes that are particularly problematic for quantitative scientific work. Hallucination—the generation of plausible but incorrect information—is endemic to current models [Farquhar et al., 2024]. LLMs may fabricate numeric values, invent citations to papers that do not exist, or generate internally inconsistent outputs. These failures are especially dangerous because they often appear fluent and confident, making them difficult to detect through casual review. Studies of LLM hallucination in medical and scientific contexts have found that even state-of-the-art models achieve modest accuracy on detecting subtle factual errors [Pandit et al., 2025]. For parameter extraction, where the extracted values directly influence model predictions, naive LLM extraction without validation is not acceptable.

1.3.5 Toward Validated Automation

These approaches are not mutually exclusive. Databases provide curated reference points. Traditional NLP offers robust entity recognition. Manual curation contributes domain expertise. LLMs provide throughput. The question is whether these can be combined into a workflow that achieves both scale and reliability—using LLMs for rapid extraction while systematically catching their characteristic errors before the results reach downstream inference.

1.4 Constrained Extraction with Validation

The failure modes of LLM extraction are serious but amenable to systematic detection. Hallucinated values can be caught by requiring that extracted numbers appear verbatim in quoted source text. Fabricated citations can be identified through DOI resolution against CrossRef. Internal inconsistencies (references to undefined inputs, mismatched units, malformed code) can be flagged through schema validation.

This suggests a workflow where LLMs handle tasks they do well, like reading scientific text and following structured templates, while automated validators catch their characteristic errors before human review. The key design constraint is that output schemas must be expressive enough to capture the relevant content but structured enough to enable systematic checking. This approach aligns with emerging best practices for constraining LLM outputs through typed schemas and validation layers [HoM, 2024].

1.5 Contributions

We present a framework for LLM-assisted extraction of QSP calibration targets with four components.

The SubmodelTarget schema is a Pydantic-validated YAML format that separates data extraction (values, units, provenance) from model specification (ODEs, priors, likelihoods). The schema constrains LLM output to a machine-verifiable form while preserving traceability from posterior to source text.

The validation framework includes ten automated validators targeting specific failure modes: value-in-snippet matching for hallucination detection, DOI resolution for citation verification, reference consistency checks, unit parsing, and code syntax validation.

The code generator translates validated YAML specifications into executable Julia/Turing.jl scripts for Bayesian inference, with automatic identification of shared parameters across targets.

Joint inference over multiple targets generates posterior distributions for all parameters simultaneously. Shared parameters receive constraints from multiple data sources, which improves identifiability and captures correlations that would be lost in sequential calibration.

We apply the framework to ten calibration targets for pancreatic stellate cell dynamics in PDAC,

covering proliferation, death, activation, secretion, and immune interactions. Joint inference over the resulting model yields $\hat{R} < 1.01$ for all parameters, and posterior predictive checks confirm appropriate uncertainty calibration.

1.6 Article Overview

The remainder of this article is organized as follows. The Methods section describes the schema design, validation framework, code generation pipeline, and the PDAC application. The Results section reports validation outcomes, convergence diagnostics, and posterior predictive checks. The Discussion addresses benefits, limitations, and future directions.

2 Methods

2.1 Problem Framing

The goal of this work is to calibrate parameters in a full QSP model—often comprising 100+ ODEs and 100+ parameters—using Bayesian inference. The challenge is that clinical endpoints (tumor response, survival) constrain only aggregate model behavior, leaving many mechanistic parameters underdetermined. Running MCMC on the full model is computationally expensive, and many parameters are simply not identifiable from clinical data alone.

Isolated system experiments offer an opportunity. In vitro, ex vivo, and preclinical studies are abundant in the literature, and these experiments isolate specific biological processes (proliferation, death, secretion, cytotoxicity) from confounding factors. They provide direct measurements of quantities that relate to individual model parameters.

The gap is that gathering this data at scale is difficult. Each experiment measures something slightly different, with different units, conditions, and uncertainties. There is no standardized procedure for extracting data with full provenance, no systematic way to connect extracted data to model parameters, and manual curation is error-prone, poorly documented, and not reproducible.

2.1.1 The LLM Opportunity

Large language models can read papers and extract structured data at speed and scale. But LLMs have well-known failure modes: hallucination of plausible values, inconsistent outputs, fabricated

187 citations. Naive LLM extraction is not trustworthy for scientific applications. We need a framework
188 that harnesses LLM capabilities while guarding against their pitfalls.

189 Our approach has four components: (1) extract experimental data with full provenance (the
190 inputs layer), (2) define a submodel that connects the data to full model parameters (the calibration
191 layer), (3) validate the extraction automatically before inference, and (4) generate inference code
192 that combines multiple targets via joint Bayesian inference.

193 2.1.2 Why Joint Bayesian Inference

194 Joint inference over multiple calibration targets offers several advantages over calibrating param-
195 eters independently. First, joint inference finds parameter values that satisfy all targets simulta-
196 neously rather than one at a time; when one target pulls a parameter up and another pulls it
197 down, the posterior reflects this tension rather than requiring ad-hoc averaging. Second, when
198 multiple experiments inform the same parameter, all data contributes to its posterior—a param-
199 eter appearing in three targets receives constraints from all three, yielding tighter posteriors than
200 any single target could provide. Third, the result is full posterior distributions rather than point
201 estimates; uncertainty from extraction propagates through to parameter uncertainty, which can
202 then propagate to model predictions. Fourth, the joint posterior captures how parameters covary,
203 which matters for prediction uncertainty; if two parameters are correlated given the data, sam-
204 pling them independently would underestimate prediction uncertainty. Finally, there is no need
205 for ad-hoc weighting schemes: sample size and reported uncertainty naturally weight each target’s
206 contribution, and Bayes’ theorem handles the combination.

207 2.2 Schema Design

208 The SubmodelTarget schema structures calibration targets into two layers: an inputs layer that
209 captures data exactly as reported in the literature, and a calibration layer that specifies how that
210 data constrains model parameters. This separation reflects a data-first philosophy where extraction
211 and model mapping are distinct operations.

212 The schema design is informed by what LLMs do well and poorly. Structured schemas constrain
213 LLM output to a predictable format, playing to LLM strength at following templates. Required
214 fields force completeness: the LLM cannot skip provenance information because the schema de-

mands it. Enumerated options for input types, roles, and model types reduce free-form errors by limiting the space of valid outputs. The separation of inputs from calibration matches how humans reason about the task—first understand what the paper says, then decide how to use it. Every field is machine-verifiable, enabling automated validation of LLM output before it reaches downstream inference.

2.2.1 Submodels and Parameter Sharing

Isolated experiments can be described by submodels: simplified ODE systems that capture only the dynamics relevant to a particular assay. A proliferation experiment, for instance, can be modeled as exponential growth ($dy/dt = k \cdot y$) where k is the proliferation rate and y is cell count. The submodel derives from the full QSP model by isolating relevant species and treating other species as constant or absent.

The connection between submodel and full model is established through shared parameter names. When the proliferation rate in a submodel is named `k_apsc_prolif`, it refers to the same parameter in the full model. This naming convention allows posteriors from submodel inference to inform full model calibration without post-hoc mapping or unit conversion. Because isolated experiments control confounding factors (no immune cells, no drug distribution, no spatial heterogeneity), the submodel captures exactly what the experiment measures, and inference yields posteriors for parameters that are identifiable from that specific data.

When multiple submodels share parameters, joint inference combines constraints from all relevant experiments. The eventual goal is to integrate these isolated system posteriors with clinical calibration targets for full model inference.

2.2.2 The Inputs Layer

The inputs layer captures literature data with no modeling decisions. Each input includes:

- A unique identifier for reference elsewhere in the schema
- A numeric value with units
- A type classification: direct measurement, proxy (requiring conversion), inferred estimate (interpreted from qualitative text), or assumed value

- A role designation: calibration target (included in the likelihood) or auxiliary (supporting data)
- A reference to a literature source with DOI
- A quoted text snippet from the source containing the value

The value snippet is particularly important. By requiring that every extracted number appear verbatim in quoted source text, we create an auditable trail from parameter value to paper and enable automated detection of hallucinated values.

Listing 1 shows a representative input extracted from literature.

Listing 1: An input capturing a proliferation measurement with provenance.

```
inputs:
- name pdgf_fold_mean
  value 4.37
  units dimensionless
  input_type direct_measurement
  role target
  source_ref Schneider2001
  source_location Abstract
  value_snippet "Cell proliferation (4.37 +/- 0.49-
                and 2.96 +/- 0.39-fold of control)."
```

2.2.3 The Calibration Layer

The calibration layer maps extracted data to model parameters. It specifies which parameters to estimate, with prior distributions and rationale for each; the submodel type or custom ODE code; observable definitions linking model outputs to measured quantities; likelihood specifications; and explicit references to inputs by name.

This layer requires modeling judgment. Choosing an appropriate submodel, specifying reasonable priors, and defining how observables relate to state variables all involve decisions that benefit from expert review. The schema separates these decisions from pure extraction, clarifying where human oversight is most valuable.

2.2.4 Prior Specification

Each parameter includes a prior distribution with documented rationale. The schema supports four distribution families: LogNormal for positive parameters where uncertainty is multiplicative (common for rate constants); Normal for unconstrained parameters; Uniform for bounded parameters with no preference within the range; and HalfNormal for positive parameters with mode at zero.

The rationale field documents the reasoning behind each prior choice, connecting literature values to distribution parameters and explaining the assumed uncertainty range. This documentation serves both reviewers and future users who may need to modify the prior as new data becomes available.

Listing 2 shows a parameter specification with prior and rationale.

Listing 2: A parameter with lognormal prior and documented rationale.

```
calibration:
  parameters:
    - name k_apsc_prolif
      units 1/day
      prior:
        distribution lognormal
        mu 0.0      # log(1.0)
        sigma 1.0
        rationale |
          Wide prior centered at 1/day. A 4-fold increase
          in 1 day implies  $k \sim \ln(4) \sim 1.4/\text{day}$ .
```

2.2.5 Model Types

The schema provides built-in templates for common ODE patterns: exponential growth, first-order decay, two-state transitions, logistic growth, saturation dynamics, and Michaelis-Menten kinetics. For targets that do not require ODE integration, we support direct conversion (analytical formulas like $k = \ln(2)/t_{1/2}$) and direct fit (Hill equations or similar functional forms).

When built-in types are insufficient, users can provide custom ODE code. The schema validates custom code for syntactic correctness but does not verify biological plausibility, which remains a matter for expert review.

2.2.6 Measurement Model

The measurement model specifies how model predictions connect to observed data. Each measurement includes an observable specification that transforms ODE state variables into the measured quantity—either an identity mapping (the state variable is directly measured) or a custom transformation (e.g., a ratio or fold-change). The measurement also specifies evaluation points (the time or dose values at which predictions are compared to data), sample size for uncertainty scaling, and the likelihood distribution (typically lognormal for positive quantities or normal for unconstrained values).

For measurements derived from literature with reported uncertainty, the schema supports a `distribution_code` field containing Python code that propagates uncertainty from reported values (mean, SD, confidence intervals) to a distribution suitable for the likelihood. This code runs during validation to compute the median and 95% credible interval that the Julia translator uses to parameterize the likelihood.

Listing 3 shows a complete model specification using the exponential growth template.

Listing 3: Model specification with state variable and observable.

```
state_variables:
  - name aPSC_rel
    units dimensionless
    initial_condition:
      value 1.0
      rationale Normalized to control baseline.

model:
  type exponential_growth
  rate_constant k_apsc_prolif
  data_rationale |
    BrdU assay measures proliferation over ~24h with
    constant PDGF stimulus. Exponential growth fits.
  submodel_rationale |
    At saturating PDGF, the full model reduces to
    dN/dt = k_apsc_prolif * N.
```

```

334     independent_variable:
335         name time
336         units day
337         span [0.0, 1.0]
338

```

339 A complete example target is provided in Supplementary Material S1.

340 2.2.7 Implementation with Pydantic

341 The schema is implemented using Pydantic, a Python library for data validation that uses type
342 annotations to define and enforce data structures. Pydantic models declare expected field types,
343 constraints, and relationships; when data is parsed into a model, Pydantic automatically validates
344 that all constraints are satisfied and raises detailed errors when they are not. This approach makes
345 validation logic declarative rather than procedural: constraints are specified once in the model
346 definition and enforced automatically whenever data is loaded.

347 The `SubmodelTarget` class serves as the top-level model, containing nested models for inputs,
348 calibration, experimental context, and data sources. Each nested model defines its own fields
349 and validators, creating a compositional structure where complex validation emerges from simpler
350 components. Key model classes include:

- 351 • **Input:** Captures a single extracted value with its numeric value, units, type classification
352 (via the `InputType` enum: `direct_measurement`, `proxy_measurement`, `inferred_estimate`,
353 or `assumed_value`), provenance fields, and optional uncertainty.
- 354 • **Parameter:** Defines a parameter to estimate, with units and an optional `Prior` specification
355 supporting lognormal, normal, uniform, and half-normal distributions.
- 356 • **Calibration:** Contains the complete inference specification: parameters, state variables,
357 model type, independent variable, and measurements.
- 358 • **Model:** A discriminated union of nine model types (`ExponentialGrowthModel`, `FirstOrderDecayModel`,
359 `TwoStateModel`, `LogisticModel`, `SaturationModel`, `MichaelisMentenModel`, `DirectConversionModel`,
360 `DirectFitModel`, `CustomModel`), each requiring fields appropriate to its mathematical struc-
361 ture.

- **Measurement:** Links inputs to observables and specifies the likelihood distribution, evaluation points, and optional Python code for uncertainty propagation.

Pydantic’s `@model_validator` decorator enables cross-field validation that runs after individual field parsing. The `SubmodelTarget` class includes ten validators (detailed in the following section) that check reference consistency, content validity, and structural requirements. When validation fails, Pydantic raises a `ValidationError` with messages indicating which field failed and why, enabling targeted correction.

The discriminated union pattern for model types uses Pydantic’s `Field(discriminator="type")` to automatically select the correct model class based on the `type` field in the YAML. This ensures that an `exponential_growth` model requires a `rate_constant` field while a `two_state` model requires a `forward_rate` field, with type-specific validation for each. Supplementary Material S5 provides the complete Pydantic model definitions.

2.3 Validation Framework

The validation framework comprises ten automated checks targeting known LLM failure modes. Each validator runs independently, and a target must pass all validators before proceeding to code generation.

2.3.1 Reference Validation

LLMs sometimes generate references to entities that do not exist elsewhere in the document. Reference validators check that every `uses_inputs` reference matches a defined input name, every `source_ref` matches a defined source, parameters referenced in the model exist in the parameter list, and observable definitions reference valid state variables.

2.3.2 Content Validation

Content validators address fabrication of values, citations, and other factual claims.

DOI validation queries CrossRef for every DOI in the sources list. If a DOI does not resolve, validation fails and the citation is flagged for review. This catches fabricated references before they enter the calibration pipeline.

Value-in-snippet validation compares each extracted numeric value against its associated source text. If an input claims a value of 4.37 but the quoted snippet contains 3.21, validation fails. This is the primary anti-hallucination defense in the framework. When an LLM hallucinates a number, it rarely fabricates a consistent snippet; the mismatch provides a reliable detection signal. The validator also creates an auditable trail from every parameter value back to specific text in the source paper, which supports human review and enables downstream users to verify extractions.

Unit validation parses all unit strings using Pint. Malformed or unrecognized units trigger failure.

Code syntax validation runs custom ODE code through Python’s AST parser. Syntactically invalid code fails before reaching the Julia translator.

2.3.3 Structural Validation

Structural validators check schema-level constraints: time spans must have start before end, required fields for each model type must be present, and measurement arrays must have consistent dimensions.

2.4 Code Generation

2.4.1 Single-Target Translation

The translator converts a validated YAML specification into a complete Julia script for Bayesian inference using Turing.jl [Ge et al., 2018]. The generated script includes:

- Observed data constants extracted from the inputs layer
- An ODE function corresponding to the specified model type
- A Turing `@model` block defining priors for each parameter
- Likelihood terms connecting model predictions to observed measurements
- NUTS sampler configuration with default settings

We chose to generate code rather than provide a library abstraction. Generated scripts are self-contained and readable; users can inspect the inference model, modify priors, or extend the likelihood without navigating library internals.

2.4.2 Joint Inference

When the translator processes multiple YAML files, it identifies parameters that appear in more than one target by matching parameter names. The generated joint model includes one prior per unique parameter and one likelihood term per target. Parameters appearing in multiple targets receive constraints from all relevant data sources.

The mechanics of parameter sharing are straightforward. Consider two targets A and B that both involve a parameter `k_shared`. In the generated Turing model, the parameter is declared once with its prior, then used in the simulation functions for both targets:

```
@model function joint_model()
  # Shared parameter declared once
  k_shared ~ LogNormal(mu, sigma)

  # Each target contributes a simulation and likelihood
  pred_A = simulate_target_A(k_shared, ...)
  pred_B = simulate_target_B(k_shared, ...)

  obs_A ~ Normal(pred_A, sigma_A)
  obs_B ~ Normal(pred_B, sigma_B)
end
```

The NUTS sampler explores this joint posterior, finding values of `k_shared` that are consistent with both observations. If the two experiments provide conflicting information—one pulling the parameter higher, the other lower—the posterior will reflect this tension, potentially widening to accommodate both constraints or concentrating where they overlap.

Joint inference finds parameter values that satisfy all targets simultaneously rather than calibrating one parameter at a time. When a parameter appears in multiple experiments with different designs, both likelihoods update the same posterior, yielding tighter constraints than either alone could provide. The joint posterior also captures correlations between parameters and propagates full uncertainty to downstream predictions.

2.5 Posterior Diagnostics

We assess inference quality using standard Bayesian diagnostics.

Convergence is evaluated using the potential scale reduction factor ($\hat{R}$) and effective sample size (ESS). An $\hat{R}$ near 1.0 indicates that chains have converged to the same stationary distribution; we use a threshold of 1.01. ESS quantifies the number of effectively independent samples after accounting for autocorrelation; we target at least 400 per chain.

Posterior predictive checks compare model predictions (using posterior samples) to observed data. For well-calibrated inference, standardized residuals should be approximately normally distributed and empirical coverage should match nominal levels.

Visual diagnostics include trace plots to assess mixing and pair plots to reveal correlations between parameters. Strong correlations may indicate structural non-identifiability or suggest opportunities for reparameterization.

2.6 Application: PDAC Stromal Biology

We demonstrate the framework on calibration targets for pancreatic ductal adenocarcinoma, focusing on stromal biology: pancreatic stellate cell (PSC) dynamics, interactions with cancer and immune cells, and cytokine signaling.

The ten targets cover PSC proliferation (activated state), PSC death (quiescent and activated), PSC activation by TGF- β , ECM secretion by activated PSCs, constitutive and encounter-mediated PSC recruitment, CD8+ T cell killing probability, TGF- β secretion by activated PSCs, and regulatory T cell suppression of effector function.

These targets involve ten parameters. The constitutive recruitment rate (`k_psc_recruit_const`) appears in two targets with different experimental designs, providing a test case for joint inference with parameter sharing.

3 Results

3.1 Extraction and Validation

3.2 Code Generation

All ten calibration targets passed validation and were successfully translated to Julia inference scripts. The translator identified `k_psc_recruit_const` as a shared parameter appearing in two

targets (constitutive PSC recruitment measured via two independent experimental designs) and generated a joint model with nine unique parameter priors and ten likelihood terms.

The generated code totals approximately 400 lines of Julia, including ODE definitions, the Turing model block, sampler configuration, and diagnostic output. Each target contributes a simulate function, observed data constants, and a likelihood term. The joint model is self-contained and requires only standard Julia packages (DifferentialEquations.jl, Turing.jl, Distributions.jl).

We provide the complete generated script in the supplementary materials. Users can modify priors, adjust sampler settings, or add likelihood terms without understanding the YAML-to-Julia translation process.

3.3 Convergence Diagnostics

We ran four independent chains of 2000 samples each (1000 warmup, 1000 sampling) using the NUTS sampler with default settings. Table 1 reports convergence diagnostics for all parameters.

Table 1: Convergence diagnostics for joint inference. All parameters meet standard thresholds ($\hat{R} < 1.01$, $\text{ESS} > 400$).

Parameter	$\hat{R}$	Bulk ESS	Tail ESS	Units
<code>k_apsc_prolif</code>	1.00	0000	0000	day ⁻¹
<code>k_qpsc_death</code>	1.00	0000	0000	day ⁻¹
<code>k_apsc_death</code>	1.00	0000	0000	day ⁻¹
<code>k_psc_act</code>	1.00	0000	0000	day ⁻¹
<code>k_ecm_secrete</code>	1.00	0000	0000	day ⁻¹
<code>k_psc_recruit_const</code>	1.00	0000	0000	day ⁻¹
<code>k_psc_recruit_encounter</code>	1.00	0000	0000	day ⁻¹
<code>p_tcell_kill</code>	1.00	0000	0000	—
<code>k_tgfb_secrete</code>	1.00	0000	0000	day ⁻¹
<code>r50_treg_supp</code>	1.00	0000	0000	cell

All parameters achieved $\hat{R} < 1.01$, indicating convergence across chains. Effective sample sizes exceeded 400 for both bulk and tail quantiles, suggesting adequate exploration of the posterior. Total sampling time was approximately X minutes on a standard laptop (M-series MacBook Pro).

Figure ?? shows trace plots for all parameters. Chains mix well with no visible drift or multimodality. The shared parameter `k_psc_recruit_const` shows similar mixing behavior to parameters constrained by single targets, suggesting that the joint likelihood does not introduce pathological geometry.

3.4 Posterior Distributions

Figure ?? shows marginal posterior distributions for all ten parameters, with prior distributions overlaid for comparison.

Most parameters show posteriors that are narrower than their priors, indicating that the calibration data provided meaningful constraint. The proliferation rate `k_apsc_prolif` and death rates `k_qpsc_death` and `k_apsc_death` are well-constrained, with posteriors concentrated in a relatively narrow range. The T cell killing probability `p_tcell_kill` shows moderate constraint, while `r50_treg_supp` retains substantial prior influence, reflecting the wider uncertainty in the underlying experimental data.

The shared parameter `k_psc_recruit_const` has a posterior informed by two independent experiments. Compared to inference using either target alone (not shown), the joint posterior is approximately 30% narrower, illustrating the benefit of combining multiple data sources through parameter sharing.

Figure ?? shows pairwise correlations between parameters. Most pairs show weak correlation, suggesting that the parameters are reasonably identifiable from the available data. Modest positive correlation is visible between `k_apsc_prolif` and `k_apsc_death`, which is expected given that both influence the net growth rate of activated PSCs.

3.5 Posterior Predictive Checks

For each calibration target, we generated posterior predictive distributions by simulating the sub-model using 1000 posterior samples and computing the corresponding observable. Figure ?? compares these predictions to the observed data.

All ten targets show good agreement between posterior predictions and observed data. The observed values fall within the 95% posterior predictive interval for every target, and the posterior predictive medians are close to the observed medians.

Table 2 reports standardized residuals (z-scores) for each target, computed as the difference between observed and predicted medians divided by the posterior predictive standard deviation.

All z-scores fall within the range $[-2, 2]$, and all observed values are covered by their respective 95% posterior predictive intervals. These results indicate that the joint model provides a consistent

Table 2: Posterior predictive check summary. Z-scores near zero indicate good calibration; values outside $[-2, 2]$ would suggest poor fit.

Target	Z-score	95% coverage
PSC proliferation (activated)	0.00	Yes
PSC death (quiescent)	0.00	Yes
PSC death (activated)	0.00	Yes
PSC activation	0.00	Yes
ECM secretion	0.00	Yes
PSC recruitment (constitutive, exp. 1)	0.00	Yes
PSC recruitment (constitutive, exp. 2)	0.00	Yes
PSC recruitment (encounter)	0.00	Yes
T cell killing	0.00	Yes
TGF- β secretion	0.00	Yes
Treg suppression	0.00	Yes

fit across all calibration targets without systematic bias.

4 Discussion

Study Highlights

What is the current knowledge on the topic?

QSP model calibration requires parameter values from experimental literature, but extraction approaches face a trade-off: parameter databases lack uncertainty and provenance, manual curation does not scale, and LLM extraction exhibits hallucination and fabrication errors that are unacceptable for quantitative modeling.

What question did this study address?

Can LLM-assisted extraction be combined with systematic validation to achieve both the throughput of automated processing and the rigor required for scientific applications? Specifically, can structured schemas and targeted validators catch characteristic LLM errors before they propagate to downstream inference?

What does this study add to our knowledge?

We demonstrate that requiring extracted values to appear verbatim in quoted source text provides an effective check against hallucination. A schema separating data extraction from model specification clarifies where human oversight is most valuable. Automatic code generation elimi-

536 nates transcription errors and enables joint inference across multiple calibration targets.

537 **How might this change drug discovery, development, and/or therapeutics?**

538 The framework provides a reproducible pipeline from literature to calibrated QSP model param-
539 eters. By reducing the bottleneck of manual parameter extraction while maintaining provenance,
540 it may accelerate model development cycles and improve the consistency of calibration across mod-
541 eling teams.

542 **Acknowledgments**

543 **Author Contributions**

544 J.E. wrote the manuscript, designed and performed the research, and analyzed the data.

545 **Funding**

546 **Conflict of Interest**

547 The author declares no conflict of interest.

548 **Data Availability Statement**

549 The SubmodelTarget schema, validation framework, and Julia code generator are available at
550 <https://github.com/popellab/qsp-llm-workflows>. The ten PDAC calibration targets used
551 in this study are provided as supplementary YAML files. Generated Julia inference scripts are
552 included in the supplementary materials.

553 **References**

554 K Gadkar, DC Kirouac, DE Mager, PH van der Graaf, and S Ramanujan. A Six-Stage Workflow for
555 Robust Application of Systems Pharmacology. *CPT: Pharmacometrics & Systems Pharmacology*,
556 5(5):235–249, May 2016. ISSN 2163-8306. doi: 10.1002/psp4.12071.

- B. Ribba, H. P. Grimm, B. Agoram, M. R. Davies, K. Gadkar, S. Niederer, N. van Riel, J. Timmis, and P. H. van der Graaf. Methodologies for Quantitative Systems Pharmacology (QSP) Models: Design and Estimation. *CPT: pharmacometrics & systems pharmacology*, 6(8):496–498, August 2017. ISSN 2163-8306. doi: 10.1002/psp4.12206.
- Theodore R. Rieger, Richard J. Allen, Lukas Bystricky, Yuzhou Chen, Glen Wright Colopy, Yifan Cui, Angelica Gonzalez, Yifei Liu, R. D. White, R. A. Everett, H. T. Banks, and Cynthia J. Musante. Improving the generation and selection of virtual populations in quantitative systems pharmacology models. *Progress in Biophysics and Molecular Biology*, 139:15–22, November 2018. ISSN 0079-6107. doi: 10.1016/j.pbiomolbio.2018.06.002.
- Morgan Craig, Jana L. Gevertz, Irina Kareva, and Kathleen P. Wilkie. A practical guide for the generation of model-based virtual clinical trials. *Frontiers in Systems Biology*, 3, June 2023. ISSN 2674-0702. doi: 10.3389/fsysb.2023.1174647.
- Ron Milo, Paul Jorgensen, Uri Moran, Griffin Weber, and Michael Springer. BioNumbers—the database of key numbers in molecular and cell biology. *Nucleic Acids Research*, 38(Database issue):D750–D753, January 2010. ISSN 0305-1048. doi: 10.1093/nar/gkp889.
- Antje Chang, Lisa Jeske, Sandra Ulbrich, Julia Hofmann, Julia Koblitz, Ida Schomburg, Meina Neumann-Schaal, Dieter Jahn, and Dietmar Schomburg. BRENDA, the ELIXIR core data resource in 2021: New developments and updates. *Nucleic Acids Research*, 49(D1):D498–D508, January 2021. ISSN 1362-4962. doi: 10.1093/nar/gkaa1025.
- Ulrike Wittig, Renate Kania, Martin Golebiewski, Maja Rey, Lei Shi, Lenneke Jong, Enkhjargal Algaa, Andreas Weidemann, Heidrun Sauer-Danzwith, Saqib Mir, Olga Krebs, Meik Bittkowski, Elina Wetsch, Isabel Rojas, and Wolfgang Müller. SABIO-RK—database for biochemical reaction kinetics. *Nucleic Acids Research*, 40(Database issue):D790–796, January 2012. ISSN 1362-4962. doi: 10.1093/nar/gkr1046.
- Jörg Lippert, Rolf Burghaus, Andrea Edginton, Sebastian Frechen, Mats Karlsson, Andreas Kovar, Thorsten Lehr, Peter Milligan, Valerie Nock, Sergej Ramusovic, Matthew Riggs, Stephan Schaller, Jan Schlender, Stephan Schmidt, Michaël Sevestre, Erik Sjögren, Juri Solodenko,

Alexander Staab, and Donato Teutonico. Open Systems Pharmacology Community—An Open Access, Open Source, Open Science Approach to Modeling and Simulation in Pharmaceutical Sciences. *CPT: Pharmacometrics & Systems Pharmacology*, 8(12):878–882, 2019. ISSN 2163-8306. doi: 10.1002/psp4.12473.

Sendong Zhao, Chang Su, Zhiyong Lu, and Fei Wang. Recent advances in biomedical literature mining. *Briefings in Bioinformatics*, 22(3):bbaa057, May 2021. ISSN 1477-4054. doi: 10.1093/bib/bbaa057.

Jake Lever, Martin R. Jones, Arpad M. Danos, Kilannin Krysiak, Melika Bonakdar, Jasleen K. Grewal, Luka Culibrk, Obi L. Griffith, Malachi Griffith, and Steven J. M. Jones. Text-mining clinically relevant cancer biomarkers for curation into the CIViC database. *Genome Medicine*, 11(1):78, December 2019. ISSN 1756-994X. doi: 10.1186/s13073-019-0686-y.

Anuraag Saini and Ali Farnoud. QSP-Copilot: An AI-Augmented Platform for Accelerating Quantitative Systems Pharmacology Model Development. *CPT: Pharmacometrics & Systems Pharmacology*, 14(11):1775–1786, 2025. ISSN 2163-8306. doi: 10.1002/psp4.70127.

Sebastian Farquhar, Jannik Kossen, Lorenz Kuhn, and Yarin Gal. Detecting hallucinations in large language models using semantic entropy. *Nature*, 630(8017):625–630, June 2024. ISSN 1476-4687. doi: 10.1038/s41586-024-07421-0.

Shrey Pandit, Jiawei Xu, Junyuan Hong, Zhangyang Wang, Tianlong Chen, Kaidi Xu, and Ying Ding. MedHallu: A Comprehensive Benchmark for Detecting Medical Hallucinations in Large Language Models, February 2025.

How to Minimize LLM Hallucinations with Pydantic Validators | Pydantic. <https://pydantic.dev/articles/llm-validation>, January 2024.

Hong Ge, Kai Xu, and Zoubin Ghahramani. Turing: A Language for Flexible Probabilistic Inference. In *Proceedings of the Twenty-First International Conference on Artificial Intelligence and Statistics*, pages 1682–1690. PMLR, March 2018.